You submitted this quiz on **Mon 28 Apr 2014 3:00 AM IST**. You got a score of **3.55** out of **5.00**. You can attempt again in 10 minutes.

Question 1

You are working on a spam classification system using regularized logistic regression. "Spam" is the positive class (y = 1) and "not spam" is the negative class (y = 0). You have trained your classifier, and there are m = 1000 examples in the cross-validation set. The chart of predicted class vs. actual class is:

For reference:

- Accuracy = (true positives + true negatives) / (total examples)
- Precision = (true positives) / (true positives + false positives)
- Recall = (true positives) / (true positives + false negatives)
- F_1 score = (2 * precision * recall) / (precision + recall)

What is the classifier's recall (as a value from 0 to 1)? Enter your answer in the box below. If necessary, provide at least two values after the decimal point.

You entered:

0.85

Your Answer		Score	Explanation
0.85	~	1.00	There are 85 true positives and 15 false negatives, so recall is 85 $/$ (85 + 15) = 0.85.
Total		1.00 / 1.00	

Question 2

Suppose a massive dataset is available for training a learning algorithm. Training on a lot of data

is likely to give good performance when two of the following conditions hold true. Which are the two?

Your Answer		Score	Explanation
When we are willing to include high order polynomial features of x (such as x_1^2 , x_2^2 , x_1x_2 , etc.).	×	0.00	As we saw with neural networks, polynomial features can still be insufficient to capture the complexity of the data, especially if the features are very high-dimensional. Instead, you should use a complex model with many parameters to fit to the large training set.
We train a model that does not use regularization.	~	0.25	Even with a very large dataset, some regularization is still likely to help the algorithm's performance, so you should use cross-validation to select the appropriate regularization parameter.
The features x contain sufficient information to predict y accurately. (For example, one way to verify this is if a human expert on the domain can confidently predict y when given only x).	×	0.00	It is important that the features contain sufficient information, as otherwise no amount of data can solve a learning problem in which the features do not contain enough information to make an accurate prediction.
✓Our learning algorithm is able to represent fairly complex functions (for example, if we train a neural network or other model with a large number of parameters).	~	0.25	You should use a complex, "low bias" algorithm, as it will be able to make use of the large dataset provided. If the mode is too simple, it will underfit the large training set.
Total		0.50 / 1.00	

Question 3

Suppose you have trained a logistic regression classifier which is outputing $h_{\theta}(x)$. Currently, you predict 1 if $h_{\theta}(x) \geq \mathrm{threshold}$, and predict 0 if $h_{\theta}(x) < \mathrm{threshold}$, where currently the threshold is set to 0.5. Suppose you **increase** the threshold to 0.9. Which of the following are true? Check all that apply.

Your Answer		Score	Explanation
The classifier is likely to have unchanged precision and recall, but higher accuracy.	~	0.25	By making more y = 0 predictions, we decrease true and false positives and increase true and false negatives. Thus, precision and recall will certainly change. We cannot say whether accuracy will increase or decrease.
The classifier is likely to now have lower recall.	~	0.25	Increasing the threshold means more y = 0 predictions. This will increase the decrease of true positives and increase the number of false negatives, so recall will decrease.
☐The classifier is likely to now have higher recall.	~	0.25	Increasing the threshold means more y = 0 predictions. This will increase the decrease of true positives and increase the number of false negatives, so recall will decrease, not increase.
☐The classifier is likely to now have lower precision.	~	0.25	Increasing the threshold means more y = 0 predictions. This will decrease both true and false positives, so precision will increase, not decrease.
Total		1.00 / 1.00	

Question 4

Suppose you are working on a spam classifier, where spam emails are positive examples (y=1) and non-spam emails are negative examples (y=0). You have a training set of emails in which 99% of the emails are non-spam and the other 1% is spam. Which of the following statements are true? Check all that apply.

Your Answer		Score	Explanation
■If you always predict non-spam	×	0.00	The classifier achieves 99% accuracy on the training set because of how skewed the classes are. We can expect

(output $y=0$), your classifier will have 99% accuracy on the training set, and it will likely perform similarly on the cross validation set.		that the cross-validation set will be skewed in the same fashion, so the classifier will have approximately the same accuracy.
If you always predict spam (output $y=1$), your classifier will have a recall of 100% and precision of 1%.	x 0.00	Since every prediction is y = 1, there are no false negatives, so recall is 100%. Furthermore, the precision will be the fraction of examples with are positive, which is 1%.
If you always predict non-spam (output $y=0$), your classifier will have 99% accuracy on the training set, but it will do much worse on the cross validation set because it has overfit the training data.	× 0.00	The classifier achieves 99% accuracy because of the skewed classes in the data, not because it is overfitting the training set. Thus, it is likely to perform just as well on the cross validation set.
If you always predict non-spam (output $y=0$), your classifier will have a recall of 0%.	✔ 0.25	Since every prediction is $y = 0$, there will be no true positives, so recall is 0% .
Total	0.25 / 1.00	

Question 5

Which of the following statements are true? Check all that apply.

Your Answer		Score	Explanation
■ After training a logistic regression	~	0.20	You can and should adjust the threshold in logistic regression using cross validation data.

classifier, you must use 0.5 as your threshold for predicting whether an example is positive or negative.			
It is a good idea to spend a lot of time collecting a large amount of data before building your first version of a learning algorithm.	~	0.20	You cannot know whether a huge dataset will be important until you have built a first version and find that the algorithm has high variance.
If your model is underfitting the training set, then obtaining more data is likely to help.	×	0.00	If the model is underfitting the training data, it has not captured the information in the examples you already have. Adding further examples will not help any more.
The "error analysis" process of manually examining the examples which your algorithm got wrong can help suggest what are good steps to take (e.g., developing new features) to improve your algorithm's performance.	~	0.20	This process of error analysis is crucial in developing high performance learning systems, as the space of possible improvements to your system is very large, and it gives you direction about what to work on next.
✓On skewed datasets (e.g., when there are more positive examples than negative examples), accuracy is not a good measure of performance and you should instead use F_1 score based on the precision and recall.	*	0.20	You can always achieve high accuracy on skewed datasets by predicting the most the same output (the most common one) for every input. Thus the ${\cal F}_1$ score is a better way to measure performance.
Total		0.80 / 1.00	